

BOUNDLESS AGENTS · AI + EDUCATION

Village of Shadows

Learning agentic AI from inside
a multi-agent world

Six AI agents · One human · One unscripted world

GOAI 2026 · Proposal Deck



Most agentic AI lessons stop at the box.

Prompt → tool → report

Useful, but it hides the hardest parts: private context, competing goals, memory, identity, human interrupts and consequences.



PROMPT

TOOL

REPORT

FIXED WORKFLOW




LIVING SYSTEM

Architecture becomes something you can sit inside.

Six independent AI agents and one human share the same world—but never the same information.

Each AI seat owns


MODEL
PERSONALITY
SECRET ROLE
PRIVATE MEMORY



The human is not supervising. The human is playing.

When the human must speak, vote or act, execution genuinely suspends.


AI and human actions then pass through the same identity and rule validation layer.




Every seat inhabits a different version of the same conversation.

Partial observability is not explained in a diagram—it changes what each agent believes and does.

Same round. Different evidence. Different truth.




Structure makes autonomous behavior safe, stateful and inspectable.



A complete game becomes a repeatable learning experiment.

The session begins with a hypothesis and ends with evidence the learner can inspect, explain and compare.



God Mode reveals what dialogue hides.

Public words are only one layer of an agent's behavior.

01 — ROLE + PRIVATE CONTEXT

02 — STATED RATIONALE

03 — MEMORY + TRUST SIGNALS

04 — TOOL CALL + DECISION

Debug view—not hidden chain-of-thought



Same world. Different minds.

Every seat can run a different provider,
model and personality.

- **CAUTIOUS** asks for another round
- **AGGRESSIVE** forces an early vote
- **DECEPTIVE** redirects suspicion
- **ANALYTICAL** connects prior claims

Preflight validation tests every configured model before the game begins.



This is a working learning lab, not a concept deck.

A complete game now closes the learning loop:
prediction, participation, live observability and an
evidence-backed debrief.

6

INDEPENDENT
AI SEATS


1

HUMAN INSIDE
THE GRAPH

4+

MODEL ROUTES
PLUS MOCK

LIVE PRODUCT SURFACE



EVIDENCE

human pause
private/public
memory growth
tool result
decision
next experiment

prediction · God Mode · checkpoints · Learning
Debrief

Safe enough for the classroom. Open enough to become more.

SAFE TO LEARN

- Synthetic game state**
- No student records required**
- Server-side secret filtering**
- Identity + rule validation**
- Reviewable event trace**

OPEN TO EXTEND

- Swap models and personalities**
- Run deterministic mock sessions**
- Reuse orchestration + tools**
- Adapt to negotiation, ethics,
incident response or security**

Educational simulation — not teacher replacement or automated student assessment.


The debrief turns one game into the next experiment.

Durable traces connect what the learner experienced to the agentic-AI concepts underneath it.

EXPERIENCE

EVIDENCE

NEXT



PLAY + PREDICT

Configure models · record a hypothesis · sit in the graph

LEARNING DEBRIEF

Interrupts · private context tools · memory · decisions

COMPARE + EXPORT

Controlled replay · lesson plans trust graph · evidence reports

VILLAGE OF SHADOWS

Do not just watch agents work. Sit among them.

Give agents identities, incomplete information
and competing objectives—then learn from
the world that emerges.

BOUNDLESS AGENTS · AI + EDUCATION · WORKING PROTOTYPE

