

BOUNDLESS AGENTS · AI + EDUCATION

Village of Shadows

Learning agentic AI from inside
a multi-agent world

Six AI agents · One human · One unscripted world


GOAI 2026 · Proposal Deck



Most agentic AI lessons stop at the box.

Prompt → tool → report

Useful, but it hides the hardest parts: private context, competing goals, memory, identity, human interrupts and consequences.



FIXED WORKFLOW




LIVING SYSTEM

Architecture becomes something you can sit inside.

Six independent AI agents and one human share the same world—but never the same information.

Each AI seat owns

MODEL
PERSONALITY
SECRET ROLE
PRIVATE MEMORY




The graph owns the world.
The agents own their choices.

The human is not supervising. The human is playing.

When the human must speak, vote or act, execution genuinely suspends.


AI and human actions then pass through the same identity and rule validation layer.




Every seat inhabits a different version of the same conversation.

Partial observability is not explained in a diagram—it changes what each agent believes and does.

Same round. Different evidence. Different truth.




Structure makes autonomous behavior safe, stateful and inspectable.



A complete game becomes a repeatable learning experiment.

The session begins with a hypothesis and ends with evidence the learner can inspect, explain and compare.



God Mode reveals what dialogue hides.

Public words are only one layer of an agent's behavior.

01 — ROLE + PRIVATE CONTEXT

02 — STATED RATIONALE

03 — MEMORY + TRUST SIGNALS

04 — TOOL CALL + DECISION

Debug view—not hidden chain-of-thought



MULTI-MODEL EXPERIMENT

Same evidence. Different minds.

Every seat can run a different provider,
model and personality.

- **CAUTIOUS** asks for another round
- **AGGRESSIVE** forces an early vote
- **DECEPTIVE** redirects suspicion
- **ANALYTICAL** connects prior claims

Preflight validation tests every configured model before the game begins.



This is a working product, not a concept deck.

A complete game loop already runs across browser UI, orchestration, agent minds, MCP actions, persistence and live events.

6

INDEPENDENT
AI SEATS


1

HUMAN INSIDE
THE GRAPH

4+

MODEL ROUTES
PLUS MOCK

LIVE PRODUCT SURFACE



GOD MODE

role
rationale
context
memory
tool call
decision

model preflight · SSE streaming · checkpoints ·
automated tests

Safe enough for the classroom. Open enough to become more.

SAFE TO LEARN

- Synthetic game state**
- No student records required**
- Server-side secret filtering**
- Identity + rule validation**
- Reviewable event trace**

OPEN TO EXTEND

- Swap models and personalities**
- Run deterministic mock sessions**
- Reuse orchestration + tools**
- Adapt to negotiation, ethics, incident response or security**

Educational simulation—not teacher replacement or automated student assessment.

The next layer turns play into measurable learning.

The orchestration core is complete. The roadmap now focuses on educator tools and comparative learning evidence.

NOW



PLAYABLE LAB

Game · traces · memory
model setup · observability

NEXT



LEARNING DEBRIEF

Guided reflection · trust graph
lesson plans · evidence rubric

EXPAND



COMPARATIVE STUDIO

Branching replay · cohorts
agent tournaments · exports

VILLAGE OF SHADOWS

Do not just watch agents work. Sit among them.

Give agents identities, incomplete information
and competing objectives—then learn from
the world that emerges.

BOUNDLESS AGENTS · AI + EDUCATION · WORKING PROTOTYPE

